

# DTSA 5510 Unsupervised Algorithms in Machine Learning Final Project

December 10, 2024

```
[31]: # Employee Performance Dataset: Unsupervised Learning Project
# Problem Description:
# The goal of this project is to perform an unsupervised learning analysis on
    ↳ an employee performance dataset. The dataset contains several features that
    ↳ reflect employee characteristics, and the task is to cluster employees into
    ↳ groups based on their performance. This can help identify performance
    ↳ patterns and assist with talent management strategies.

# Type of Learning: Unsupervised learning
# Task: Clustering using K-Means and Dimensionality Reduction (PCA)
```

```
[ ]: # Data Collection and Description
# The dataset used for this analysis is titled "Employee_Performance_dataset.
    ↳ csv". It was obtained from a publicly available source (insert source or
    ↳ citation). This dataset contains information about employees' attributes
    ↳ such as work experience, performance ratings, and more.

# Data Size:
# Number of rows (samples): 1,000 employees
# Number of columns (features): 10 columns (including 'EmployeeID', 'Age',
    ↳ 'Performance', 'WorkExperience', etc.)
# Data Features:
# EmployeeID: Unique identifier for each employee
# Age: Age of the employee
# Performance: Performance rating of the employee (numeric)
# WorkExperience: Years of work experience
# Other features: Additional features such as education level, department, etc.
```

```
[29]: import pandas as pd

# Load the dataset
df = pd.read_csv('Employee_Performance_dataset.csv')

# Display basic information
print(df.head()) # Display first few rows
print(df.info()) # Data types and missing values
```

```
print(df.describe()) # Summary statistics
```

	ID	Name	Age	Gender	Department	Salary	Joining Date	\
0	1	Cory Escobar	48	Female	HR	5641	2015-05-03	
1	2	Timothy Sanchez	25	Other	Sales	4249	2020-11-09	
2	3	Chad Nichols	57	Other	Sales	3058	2019-02-12	
3	4	Christine Williams	58	Female	IT	5895	2017-09-08	
4	5	Amber Harris	35	Other	IT	4317	2020-02-15	

	Performance Score	Experience	Status	Location	Session
0	2.0	16	Active	New York	Night
1	2.0	11	Inactive	Los Angeles	Evening
2	NaN	1	Inactive	New York	Morning
3	2.0	13	Inactive	Los Angeles	Evening
4	5.0	16	Inactive	New York	Evening

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1000 entries, 0 to 999
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	ID	1000 non-null	int64
1	Name	1000 non-null	object
2	Age	1000 non-null	int64
3	Gender	1000 non-null	object
4	Department	1000 non-null	object
5	Salary	1000 non-null	int64
6	Joining Date	1000 non-null	object
7	Performance Score	502 non-null	float64
8	Experience	1000 non-null	int64
9	Status	1000 non-null	object
10	Location	1000 non-null	object
11	Session	1000 non-null	object

```
dtypes: float64(1), int64(4), object(7)
```

```
memory usage: 93.9+ KB
```

```
None
```

	ID	Age	Salary	Performance Score	Experience
count	1000.000000	1000.000000	1000.000000	502.000000	1000.000000
mean	500.500000	40.782000	5917.374000	2.910359	10.120000
std	288.819436	14.124871	2299.418003	1.424736	5.713689
min	1.000000	18.000000	2015.000000	1.000000	1.000000
25%	250.750000	28.000000	3829.750000	2.000000	5.000000
50%	500.500000	40.000000	5889.000000	3.000000	10.000000
75%	750.250000	52.000000	7903.250000	4.000000	15.000000
max	1000.000000	65.000000	9993.000000	5.000000	20.000000

```
[20]: # Data Cleaning
      # Handling Missing Values:
```

```
# First, we check for missing values in the dataset:
```

```
import pandas as pd
```

```
df = pd.read_csv('Employee_Performance_dataset.csv')  
print(df.isnull().sum()) # Check for missing values
```

```
ID                0  
Name              0  
Age              0  
Gender            0  
Department        0  
Salary            0  
Joining Date      0  
Performance Score 498  
Experience         0  
Status            0  
Location          0  
Session           0  
dtype: int64
```

```
[21]: # If there are any missing values, we can handle them by imputing or removing  
      ↳ the rows/columns:
```

```
# Impute missing values with the mean of each column  
df = df.fillna(df.mean())
```

```
[30]: print(df.dtypes) # Check data types
```

```
ID                int64  
Name              object  
Age              int64  
Gender            object  
Department        object  
Salary            int64  
Joining Date      object  
Performance Score float64  
Experience         int64  
Status            object  
Location          object  
Session           object  
dtype: object
```

```
[22]: # Handling Infinite Values:  
      # We replace infinite values with NaN and handle them:
```

```
import numpy as np
```

```
df.replace([np.inf, -np.inf], np.nan, inplace=True)
df = df.fillna(df.mean()) # Impute again after replacing infinities
```






```
[23]: # Exploratory Data Analysis (EDA)
# The purpose of this EDA is to better understand the dataset before applying
↳ unsupervised learning algorithms.

# Visualizations:
# We start by visualizing the distribution of key features using histograms and
↳ box plots:

import matplotlib.pyplot as plt

# Histograms for numeric features
df.hist(bins=15, figsize=(15, 10))
plt.show()

# Box plot for performance metric
df.boxplot(column='Performance') # Adjust column name as needed
plt.show()
```



```

KeyError                                Traceback (most recent call
↳last)

<ipython-input-23-6ce19541a6a7> in <module>
    12
    13 # Box plot for performance metric
----> 14 df.boxplot(column='Performance') # Adjust column name as needed
    15 plt.show()

/opt/conda/lib/python3.7/site-packages/pandas/plotting/_core.py in
↳boxplot_frame(self, column, by, ax, fontsize, rot, grid, figsize, layout,
↳return_type, backend, **kwargs)
    445         layout=layout,
    446         return_type=return_type,
--> 447         **kwargs,
    448     )
    449

/opt/conda/lib/python3.7/site-packages/pandas/plotting/_matplotlib/
↳boxplot.py in boxplot_frame(self, column, by, ax, fontsize, rot, grid,
↳figsize, layout, return_type, **kwds)
    373         layout=layout,
    374         return_type=return_type,
--> 375         **kwds,
    376     )
    377     plt.draw_if_interactive()

/opt/conda/lib/python3.7/site-packages/pandas/plotting/_matplotlib/
↳boxplot.py in boxplot(data, column, by, ax, fontsize, rot, grid, figsize,
↳layout, return_type, **kwds)
    339         columns = data.columns
    340     else:
--> 341         data = data[columns]
    342
    343     result = plot_group(columns, data.values.T, ax)

/opt/conda/lib/python3.7/site-packages/pandas/core/frame.py in
↳__getitem__(self, key)
   2804         if is_iterator(key):
   2805             key = list(key)

```

```

-> 2806             indexer = self.loc._get_listlike_indexer(key, axis=1,
↳raise_missing=True)[1]
    2807
    2808             # take() does not accept boolean indexers

```

```

/opt/conda/lib/python3.7/site-packages/pandas/core/indexing.py in
↳_get_listlike_indexer(self, key, axis, raise_missing)
    1551
    1552         self._validate_read_indexer(
-> 1553             keyarr, indexer, o._get_axis_number(axis),
↳raise_missing=raise_missing
    1554         )
    1555         return keyarr, indexer

```

```


/opt/conda/lib/python3.7/site-packages/pandas/core/indexing.py in
↳_validate_read_indexer(self, key, indexer, axis, raise_missing)
    1638             if missing == len(indexer):
    1639                 axis_name = self.obj._get_axis_name(axis)
-> 1640                 raise KeyError(f"None of [{key}] are in the
↳[{axis_name}]")
    1641
    1642             # We (temporarily) allow for some missing keys with .
↳loc, except in

```


```

KeyError: "None of [Index(['Performance'], dtype='object')] are in the
↳[columns]"

```



```
[24]: # Correlation Matrix:  
# We explore relationships between features:  
  
import seaborn as sns  
  
plt.figure(figsize=(10, 8))  
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')  
plt.show()
```



```
[25]: # Data Transformation (Scaling)
# Since K-Means is sensitive to the scale of the data, we apply scaling to
# → normalize the numeric features:

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
df_scaled = scaler.fit_transform(df.select_dtypes(include=['float64', 'int64']))
```

```
[26]: # Model Building: K-Means Clustering
# We apply the K-Means algorithm to group employees into clusters based on
# → their performance and characteristics.

# Model Training:

from sklearn.cluster import KMeans

# Apply K-Means clustering with 3 clusters
kmeans = KMeans(n_clusters=3, random_state=42)
clusters = kmeans.fit_predict(df_scaled)
```




```
# Add cluster labels to the dataframe
df['Cluster'] = clusters
```

```
[27]: # Visualizing Clusters:
# We use PCA for dimensionality reduction to visualize the clusters in a 2D
      ↪ space:

from sklearn.decomposition import PCA

pca = PCA(n_components=2)
pca_components = pca.fit_transform(df_scaled)

# Plot the results
plt.scatter(pca_components[:, 0], pca_components[:, 1], c=clusters)
plt.title('PCA of Employee Performance Clusters')
plt.xlabel('PC1')
plt.ylabel('PC2')
plt.show()
```



```
[2]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
```

```

from sklearn.preprocessing import StandardScaler

# Load the dataset (assuming it is already available as 'df')
df = pd.read_csv('Employee_Performance_dataset.csv')

# Handle missing values in 'Performance Score'
df['Performance Score'].fillna(df['Performance Score'].mean(), inplace=True)

# Scaling the numerical features
numeric_cols = ['Age', 'Salary', 'Performance Score', 'Experience']
scaler = StandardScaler()
df_scaled = scaler.fit_transform(df[numeric_cols])

# Elbow Method to determine the optimal number of clusters
def perform_kmeans(df_scaled, max_clusters=10):
    inertia = []
    for k in range(1, max_clusters + 1):
        kmeans = KMeans(n_clusters=k, random_state=42)
        kmeans.fit(df_scaled)
        inertia.append(kmeans.inertia_)

    plt.figure(figsize=(8, 6))
    plt.plot(range(1, max_clusters + 1), inertia, marker='o')
    plt.title('Elbow Method for Optimal K')
    plt.xlabel('Number of clusters')
    plt.ylabel('Inertia (Sum of Squared Distances)')
    plt.show()

# Perform the elbow method to find the optimal K
perform_kmeans(df_scaled)

# Perform K-Means clustering with an optimal number of clusters (4 for now)
optimal_k = 4
kmeans_model = KMeans(n_clusters=optimal_k, random_state=42)
kmeans_model.fit(df_scaled)

# Add cluster labels to the dataframe
df['Cluster'] = kmeans_model.labels_

# Perform PCA for visualization
def perform_pca_and_plot(df_scaled, n_components=2):
    pca = PCA(n_components=n_components)
    pca_components = pca.fit_transform(df_scaled)
    pca_df = pd.DataFrame(data=pca_components, columns=[f'PCA{i+1}' for i in
    →range(n_components)])
    pca_df['Cluster'] = df['Cluster']

```


```


# Plot the PCA components with clusters
plt.figure(figsize=(8, 6))
scatter = plt.scatter(pca_df['PCA1'], pca_df['PCA2'], c=pca_df['Cluster'],
→ cmap='viridis')
plt.title('PCA of Employee Performance Clusters')
plt.xlabel('PCA 1')
plt.ylabel('PCA 2')
plt.colorbar(scatter)
plt.show()

# Perform PCA and visualize the clusters
perform_pca_and_plot(df_scaled)

# Print the centroids of the KMeans model
print("Cluster Centers (Centroids):")
print(kmeans_model.cluster_centers_)

```





Cluster Centers (Centroids):

```
[[-0.13032584  0.79725327 -0.04584479 -0.93169828]
 [ 1.01024503 -0.7516175  -0.02938875  0.00363182]
 [-0.03889579  0.91794178 -0.16113947  0.99044341]
 [-0.91465709 -0.81182985  0.21864408  0.09294507]]
```

```
[28]: # Results and Analysis
# After applying the K-Means algorithm, we analyze the clustering results:

# The dataset was divided into 3 clusters.
# Each cluster represents a distinct group of employees with similar
    ↳ performance characteristics.
# Evaluation Metrics:
# Since this is an unsupervised learning task, we use visualization and cluster
    ↳ analysis (e.g., the silhouette score or cluster centers) to assess the
    ↳ quality of the clustering:

from sklearn.metrics import silhouette_score
```

```
score = silhouette_score(df_scaled, clusters)
print(f'Silhouette Score: {score}')
```

Silhouette Score: 0.14591961213673368

```
[ ]: #Conclusion
#Based on the clustering results, we can conclude that the K-Means algorithm
↳ successfully grouped employees into 3 clusters based on their performance
↳ and characteristics. However, further tuning and additional unsupervised
↳ models such as DBSCAN or Hierarchical Clustering could be explored to
↳ improve the clustering accuracy.
```